

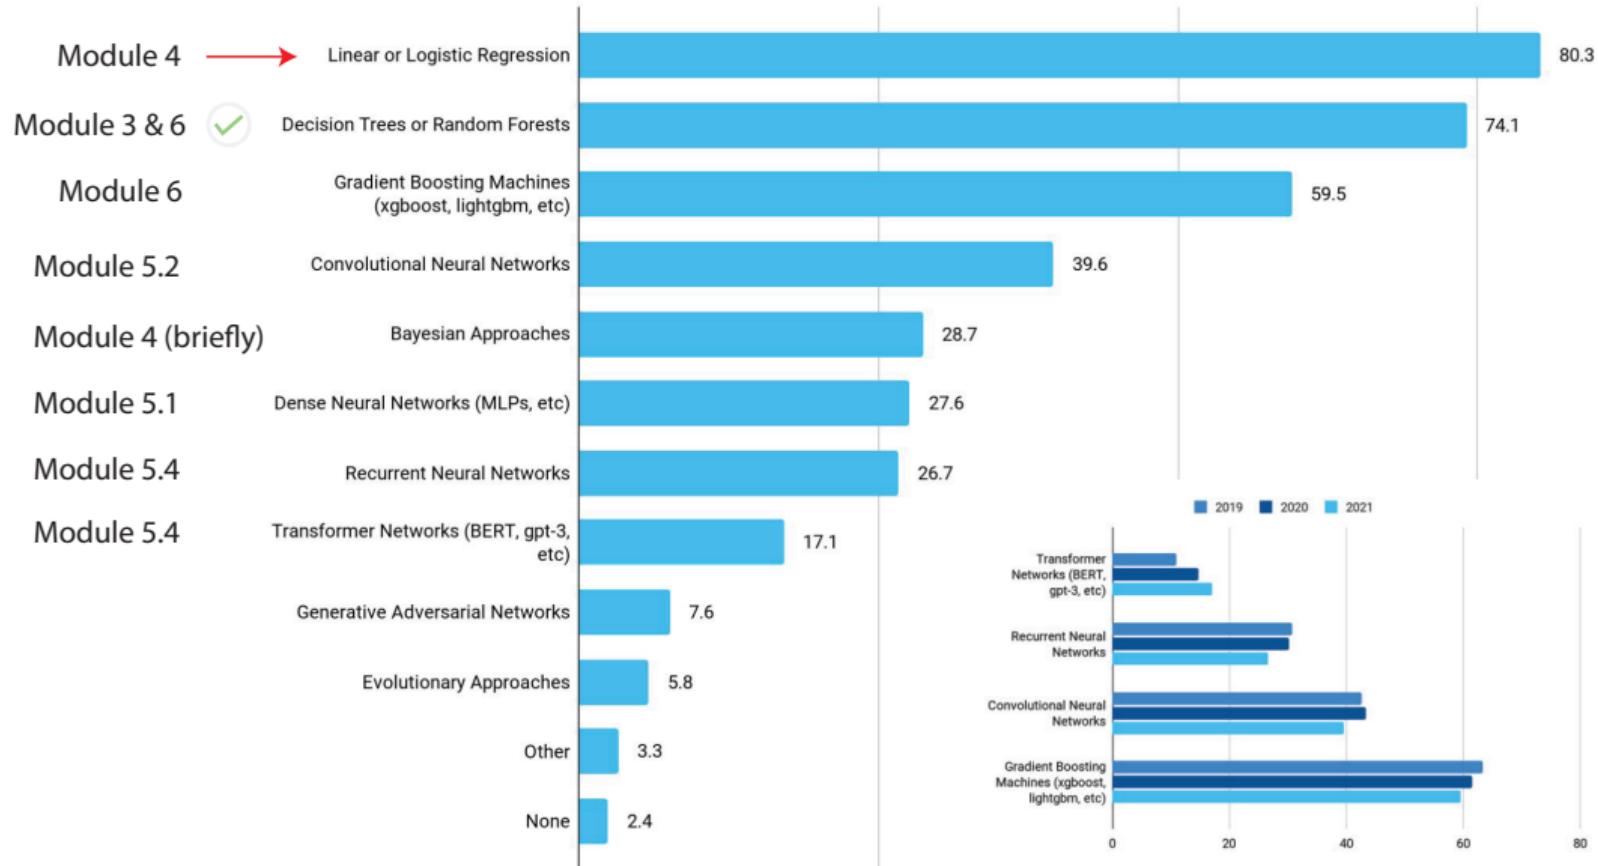
Lecture 7-9 - Module 4.1. Linear regression

COMP 551 Applied machine learning

Yue Li
Assistant Professor
School of Computer Science
McGill University

January 28, 30, & Feb 4, 2025

Most commonly used ML algorithms in Kaggle 2021 survey



Outline

Objectives

Simple linear regression

Multiple linear regression

Probabilistic interpretation

Non-linear basis function

Summary

Outline

Objectives

Simple linear regression

Multiple linear regression

Probabilistic interpretation

Non-linear basis function

Summary

Learning objectives

Understanding the following concepts

- Simple linear model
- Finding the best linear fit by minimizing SSE
- Matrix algebra in solving multiple regression
- Probabilistic interpretation
- Feature transformation by non-linear basis functions

Outline

Objectives

Simple linear regression

Multiple linear regression

Probabilistic interpretation

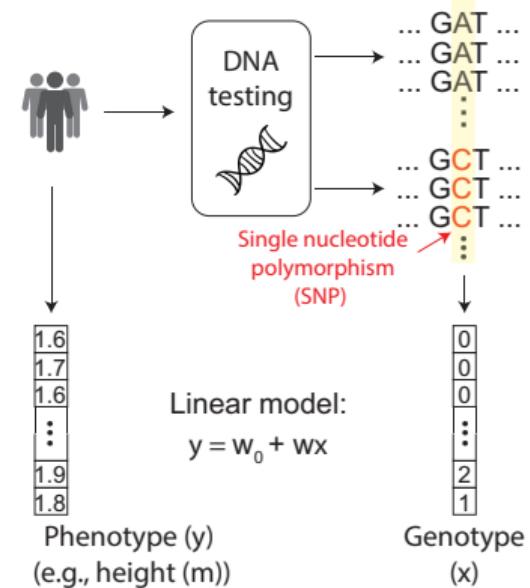
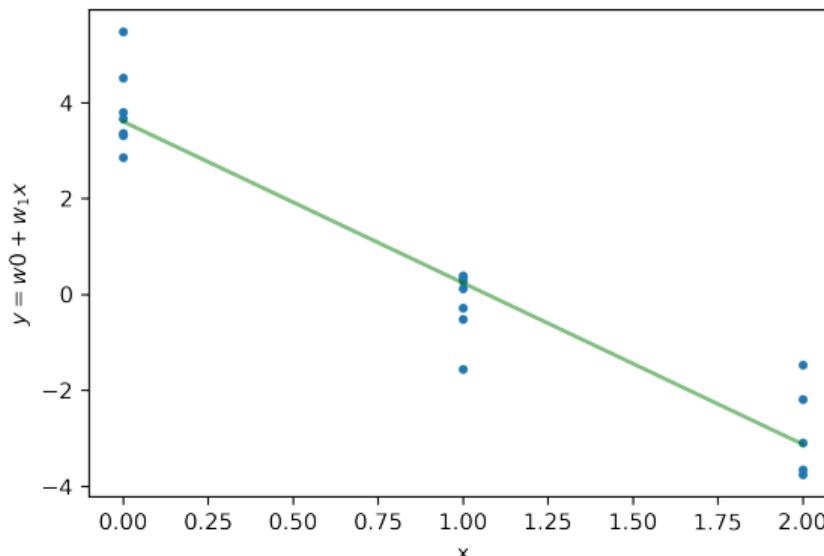
Non-linear basis function

Summary

Linear regression using a one-dimensional input

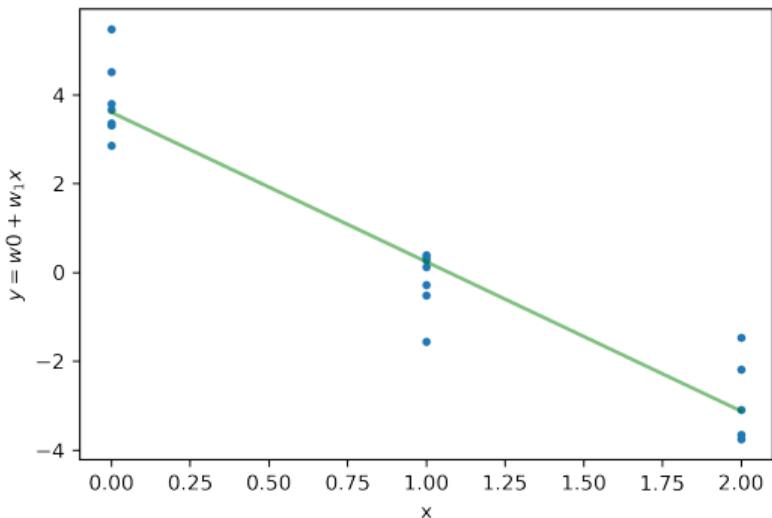
We want to predict real-valued quantity or often known as **response or target variable** $y \in \mathbb{R}$ by finding a mapping function that maps from a **one-dimensional input** x to the real-valued y . We can fit a linear function on training examples $\{x^{(n)}, y^{(n)}\}_{n=1}^N$ (e.g., predicting phenotype from one genetic mutation):

$$f(x^{(n)}; w_0, w_1) = w_0 + w_1 x^{(n)} \quad (1)$$



Linear regression using a one-dimensional input

$$f(x^{(n)}; w_0, w_1) = w_0 + w_1 x^{(n)}$$



- $w_0 = 3.4$ is the **intercept** or **bias**, which is not be confused with the “model bias”
- $w_1 = -3.25$ is the **slope** of the linear function or the **regression coefficient**.

Simple linear regression using one input feature

In statistics, we often write down the regression formula as:

$$y^{(n)} = w_0 + w_1 x^{(n)} + \epsilon^{(n)}$$

where $\epsilon^{(n)}$ is the prediction error for example n .

Using matrix notation, we can write the regression formula as:

$$\begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(N)} \end{bmatrix} = \begin{bmatrix} 1 & x^{(1)} \\ 1 & x^{(2)} \\ \vdots & \vdots \\ 1 & x^{(N)} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} + \begin{bmatrix} \epsilon^{(1)} \\ \epsilon^{(2)} \\ \vdots \\ \epsilon^{(N)} \end{bmatrix} \quad (2)$$

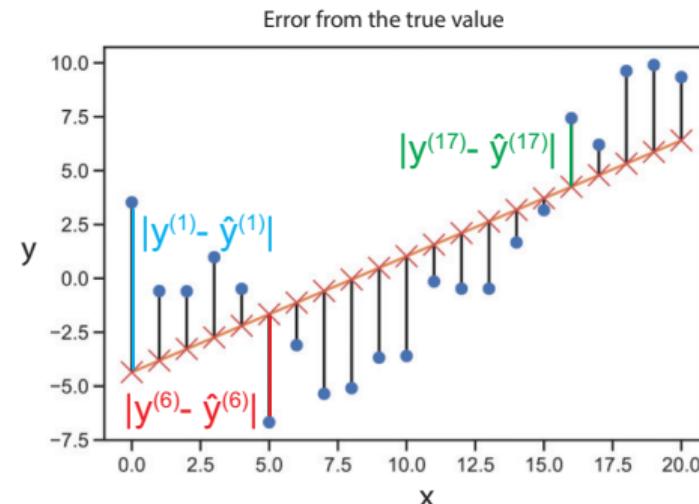
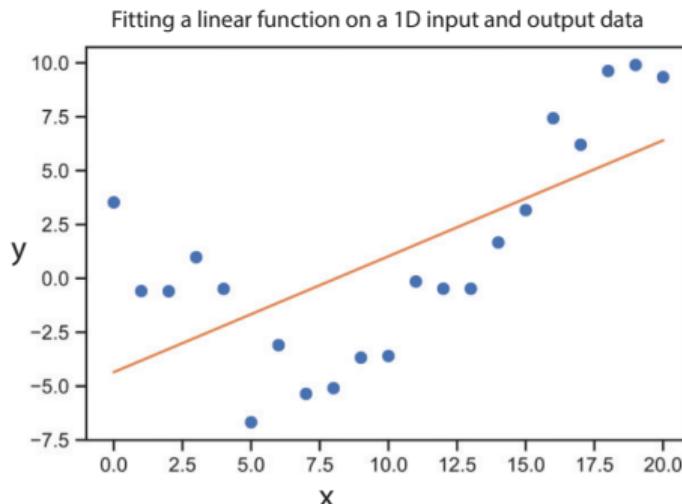
$$\mathbf{y} = \mathbf{X}\mathbf{w} + \boldsymbol{\epsilon} \quad (3)$$

where \mathbf{y} and $\boldsymbol{\epsilon}$ are both $N \times 1$ vectors, \mathbf{X} is a $N \times 2$ matrix, and \mathbf{w} is a 2×1 vector.

Residual error as a measure of prediction loss

Every straight line we fit incurs a prediction error on the training data point unless the fitted line goes through that data point. The **residual error** is Euclidean distance between the observed response $y^{(n)}$ value and the predicted response $\hat{y}^{(n)} = \mathbf{x}^{(n)}\mathbf{w}$:

$$l_n = \|y^{(n)} - \hat{y}^{(n)}\|_2 = \sqrt{(y^{(n)} - \hat{y}^{(n)})^2} = |y^{(n)} - \hat{y}^{(n)}| \quad (4)$$



Notation: $\|\mathbf{a}\|_2 = \sqrt{\sum_i a_i^2}$ is L_2 -norm and $\|\mathbf{a}\|_2^2 = \sum_i a_i^2$ is the squared of the L2-norm.

Fitting a linear regression function by minimizing the sum of squared error

Sum of squared error (SSE) as a function of the linear coefficients \mathbf{w} is defined as:

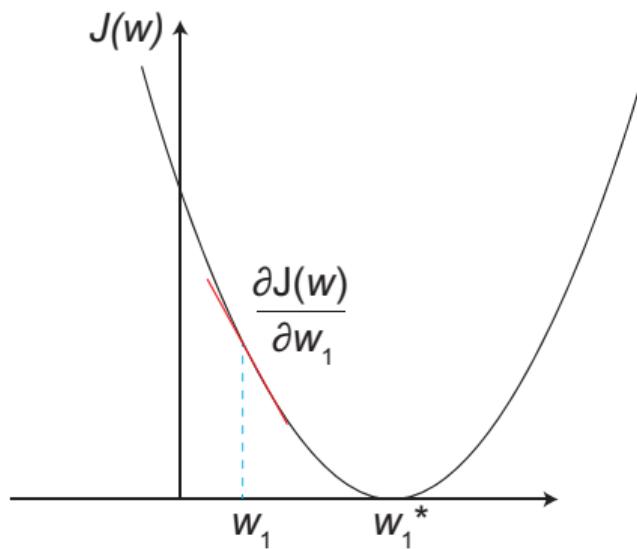
$$J(\mathbf{w}) = \|\mathbf{y} - \hat{\mathbf{y}}\|_2^2 = \sum_{n=1}^N (y^{(n)} - \hat{y}^{(n)})^2 = \sum_{n=1}^N (y^{(n)} - w_0 - w_1 x^{(n)})^2 = (\mathbf{y} - \hat{\mathbf{y}})^\top (\mathbf{y} - \hat{\mathbf{y}}) = \boldsymbol{\epsilon}^\top \boldsymbol{\epsilon} \quad (5)$$

Goal: find the best \mathbf{w} that minimizes the SSE:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} J(\mathbf{w}) \quad (6)$$

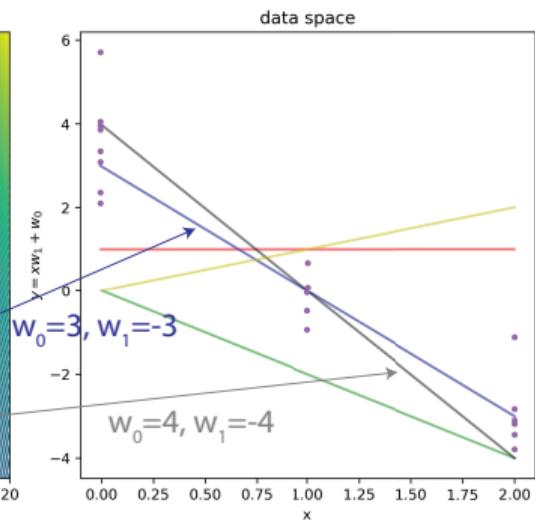
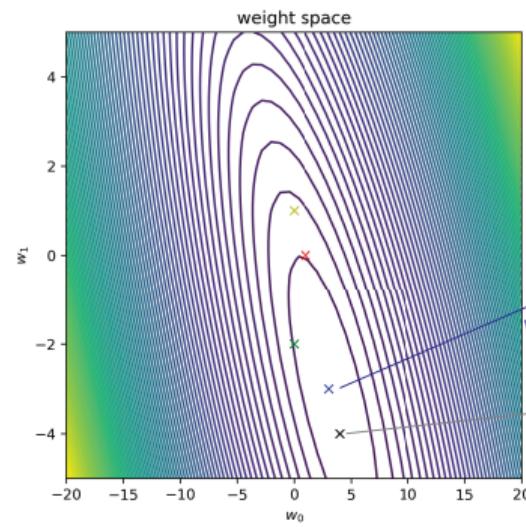
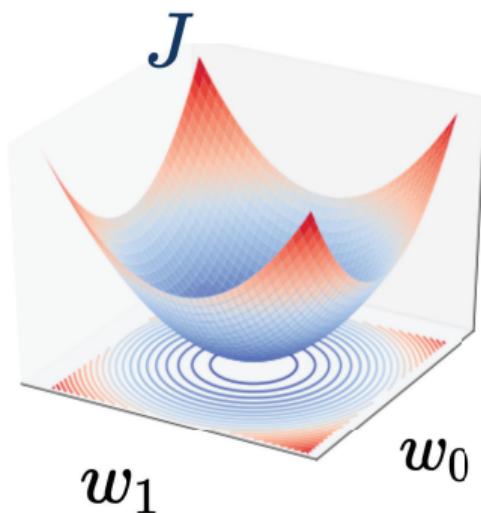
Given the error function is differentiable everywhere, the slope at the current value of w_1 projecting onto the error parabola is shown on the right.

The SSE error function is **convex** (more details in Module 4.4). The optimal w_1^* is found where **the slope or the partial derivative is zero** $\frac{\partial J(\mathbf{w})}{\partial w_1^*} = 0$.



Contour plot visualizes the error surface as a function of w_1 and w_0

$$J(\mathbf{w}) = \sum_{n=1}^N (y^{(n)} - \hat{y}^{(n)})^2 = \sum_{n=1}^N (y^{(n)} - w_0 - w_1 x^{(n)})^2$$



Derivation of ordinary least squared (OLS) solution for w_0

Solving for w_0 :

$$\begin{aligned}\frac{\partial J(\mathbf{w})}{\partial w_0} &= \frac{\partial}{\partial w_0} \sum_{n=1}^N (y^{(n)} - w_0 - w_1 x^{(n)})^2 \\ &= \sum_{n=1}^N \frac{\partial(y^{(n)} - w_0 - w_1 x^{(n)})^2}{\partial(y^{(n)} - w_0 - w_1 x^{(n)})} \frac{\partial(y^{(n)} - w_0 - w_1 x^{(n)})}{\partial w_0} \\ &= \sum_{n=1}^N 2(y^{(n)} - w_0 - w_1 x^{(n)})(-1)\end{aligned}$$

Set $\frac{\partial J(w_0)}{\partial w_0}$ to zero and multiplying $\frac{1}{2}$ and -1 on both sides gives:

$$\sum_{n=1}^N (y^{(n)} - w_0 - w_1 x^{(n)}) = 0$$

$$\begin{aligned}\sum_{n=1}^N y^{(n)} - \sum_{n=1}^N w_0 - \sum_{n=1}^N w_1 x^{(n)} &= 0 \\ \sum_{n=1}^N w_0 &= \sum_{n=1}^N y^{(n)} - w_1 \sum_{n=1}^N x^{(n)} \\ Nw_0 &= \sum_{n=1}^N y^{(n)} - w_1 \sum_{n=1}^N x^{(n)} \\ w_0 &= \frac{1}{N} \sum_{n=1}^N y^{(n)} - w_1 \frac{1}{N} \sum_{n=1}^N x^{(n)}\end{aligned}$$

$$w_0 = \bar{y} - w_1 \bar{x}$$

Therefore, $\hat{w}_0 = \bar{y} - w_1 \bar{x}$

Derivation of OLS solution for w_1

$$\begin{aligned}\frac{\partial J(\mathbf{w})}{\partial w_1} &= \frac{\partial}{\partial w_1} \sum_{n=1}^N (y^{(n)} - w_0 - w_1 x^{(n)})^2 \\ &= \sum_{n=1}^N 2(y^{(n)} - w_0 - w_1 x^{(n)})(-x^{(n)})\end{aligned}$$

Set $\frac{\partial J(\mathbf{w})}{\partial w_1}$ to zero and multiplying $\frac{1}{2}$ and -1 on both sides gives:

$$\sum_{n=1}^N (y^{(n)} - w_0 - w_1 x^{(n)})x^{(n)} = 0 \quad (7)$$

Plug the bias estimate \hat{w}_0 in (7) and solve for w_1 :

$$\sum_{n=1}^N (y^{(n)} - (\bar{y} - w_1 \bar{x}) - w_1 x^{(n)})x^{(n)} = 0$$

$$\sum_{n=1}^N (y^{(n)} - \bar{y} + w_1 \bar{x} - w_1 x^{(n)})x^{(n)} = 0$$

$$\sum_{n=1}^N (y^{(n)} - \bar{y})x^{(n)} - w_1 \sum_{n=1}^N (x^{(n)} - \bar{x})x^{(n)} = 0$$

$$\hat{w}_1 = \frac{\sum_{n=1}^N (y^{(n)} - \bar{y})x^{(n)}}{\sum_{n=1}^N (x^{(n)} - \bar{x})x^{(n)}}$$

Derivation of OLS solution for w_1 (cont'd)

Note that

$$\sum_{n=1}^N (y^{(n)} - \bar{y})x^{(n)} = \sum_{n=1}^N (y^{(n)} - \bar{y})(x^{(n)} - \bar{x})$$

because:

$$\begin{aligned}\sum_{n=1}^N (y^{(n)} - \bar{y})(x^{(n)} - \bar{x}) &= \sum_{n=1}^N y^{(n)}x^{(n)} - \sum_{n=1}^N y^{(n)}\bar{x} - \sum_{n=1}^N \bar{y}x^{(n)} + \sum_{n=1}^N \bar{y}\bar{x} \\ &= \sum_{n=1}^N y^{(n)}x^{(n)} - N\bar{y}\bar{x} - \sum_{n=1}^N \bar{y}x^{(n)} + N\bar{y}\bar{x} = \sum_{n=1}^N (y^{(n)} - \bar{y})x^{(n)}\end{aligned}$$

Similarly,

$$\sum_{n=1}^N (x^{(n)} - \bar{x})x^{(n)} = \sum_{n=1}^N (x^{(n)} - \bar{x})(x^{(n)} - \bar{x}) = \sum_{n=1}^N (x^{(n)} - \bar{x})^2$$

Update equations for simple linear regression

Therefore, the simple linear regression OLS solutions are:

$$\hat{w}_1 = \frac{\sum_{n=1}^N (y^{(n)} - \bar{y})x^{(n)}}{\sum_{n=1}^N (x^{(n)} - \bar{x})x^{(n)}} = \frac{\sum_{n=1}^N (y^{(n)} - \bar{y})(x^{(n)} - \bar{x})}{\sum_{n=1}^N (x^{(n)} - \bar{x})^2}; \quad \hat{w}_0 = \bar{y} - \hat{w}_1 \bar{x} \quad (8)$$

Compare this solution with Pearson correlation coefficient (PCC) yields useful insights:

$$\hat{r}_{xy} = \frac{\sum_{n=1}^N (y^{(n)} - \bar{y})(x^{(n)} - \bar{x})}{\sqrt{\sum_{n=1}^N (x^{(n)} - \bar{x})^2} \sqrt{\sum_{n=1}^N (y^{(n)} - \bar{y})^2}} \in [-1, 1]$$

Therefore, we can see that

$$\hat{w}_1 = \hat{r}_{xy} \frac{\sqrt{\sum_{n=1}^N (y^{(n)} - \bar{y})^2}}{\sqrt{\sum_{n=1}^N (x^{(n)} - \bar{x})^2}}$$

If the standard deviation for y and x are the same (e.g., through standardization described next), $\hat{w}_1 = \hat{r}_{xy}$. However, in general, the regression slope and PCC are *not* the same: while \hat{r}_{xy} gives a bounded measure independent of the scale of the two variables, \hat{w}_1 measures the change in the expected value of y corresponding to 1-unit increase/decrease of x .

Update equations for simple linear regression

A special case arises when \mathbf{y} and \mathbf{x} are *centered*: $\dot{\mathbf{y}}^{(n)} = \mathbf{y}^{(n)} - \bar{\mathbf{y}}$ and $\dot{\mathbf{x}}^{(n)} = \mathbf{x}^{(n)} - \bar{\mathbf{x}}$, such that $\bar{\mathbf{y}} = 0$ and $\bar{\mathbf{x}} = 0$ because

$$\bar{\dot{\mathbf{y}}} = \frac{1}{N} \sum_n \dot{\mathbf{y}}^{(n)} = \frac{1}{N} \sum_n (\mathbf{y}^{(n)} - \bar{\mathbf{y}}) = \frac{1}{N} \sum_n \mathbf{y}^{(n)} - \frac{1}{N} \sum_n \bar{\mathbf{y}} = \bar{\mathbf{y}} - \bar{\mathbf{y}} = 0$$

As a result,

$$\hat{w}_1 = \frac{\sum_{n=1}^N (\dot{\mathbf{y}}^{(n)} - \bar{\mathbf{y}})(\dot{\mathbf{x}}^{(n)} - \bar{\mathbf{x}})}{\sum_{n=1}^N (\dot{\mathbf{x}}^{(n)} - \bar{\mathbf{x}})^2} = \frac{\sum_{n=1}^N \dot{\mathbf{y}}^{(n)} \dot{\mathbf{x}}^{(n)}}{\sum_{n=1}^N (\dot{\mathbf{x}}^{(n)})^2}; \quad \hat{w}_0 = \bar{\mathbf{y}} - \hat{w}_1 \bar{\mathbf{x}} = 0$$

Without the intercept, the linear function becomes $\hat{\mathbf{y}}^{(n)} = w_1 \dot{\mathbf{x}}^{(n)}$

If we only center the input variable such that $\bar{\mathbf{x}} = 0$ but not \mathbf{y} , then $w_0 = \bar{\mathbf{y}}$ and the linear function $\hat{\mathbf{y}}^{(n)} = w_1 \dot{\mathbf{x}}^{(n)} + \bar{\mathbf{y}}$ measures the *change* from the average $\bar{\mathbf{y}}$.

Also, PCC of centered \mathbf{x} and \mathbf{y} is the same as cosine similarity between \mathbf{x} and \mathbf{y} .

Standardization involves centering and scaling the input feature

Another special case arises if we further scale \dot{x} by its standard deviation

$$\sigma_{\dot{x}} = \sqrt{\frac{1}{N} \sum_n (\dot{x}^{(n)})^2}$$

$$\tilde{\dot{x}}^{(n)} = \dot{x}^{(n)} / \sigma_{\dot{x}}$$

then we have

$$\sigma_{\tilde{\dot{x}}}^2 = \frac{1}{N} \sum_n (\tilde{\dot{x}}^{(n)})^2 = \frac{1}{N} \sum_n (\dot{x}^{(n)} / \sigma_{\dot{x}})^2 = \frac{1}{N} \sum_n (\dot{x}^{(n)})^2 / \sigma_{\dot{x}}^2 = \sigma_{\dot{x}}^2 / \sigma_{\dot{x}}^2 = 1$$

The regression coefficient becomes more simplified (while $w_0 = 0$ after centering x, y):

$$\hat{w}_1 = \frac{\sum_{n=1}^N \tilde{\dot{y}}^{(n)} \tilde{\dot{x}}^{(n)}}{\sum_{n=1}^N (\tilde{\dot{x}}^{(n)})^2} = \frac{1}{N} \sum_{n=1}^N \tilde{\dot{y}}^{(n)} \tilde{\dot{x}}^{(n)} = \frac{1}{N} \tilde{\dot{\mathbf{x}}}^T \tilde{\dot{\mathbf{y}}}$$

Together, the procedure of centering and scaling of the response and/or the input features is common practice known as **standardization** mainly to simplify computation and to make the model robust to different numerical scales of the input and response.

Outline

Objectives

Simple linear regression

Multiple linear regression

Probabilistic interpretation

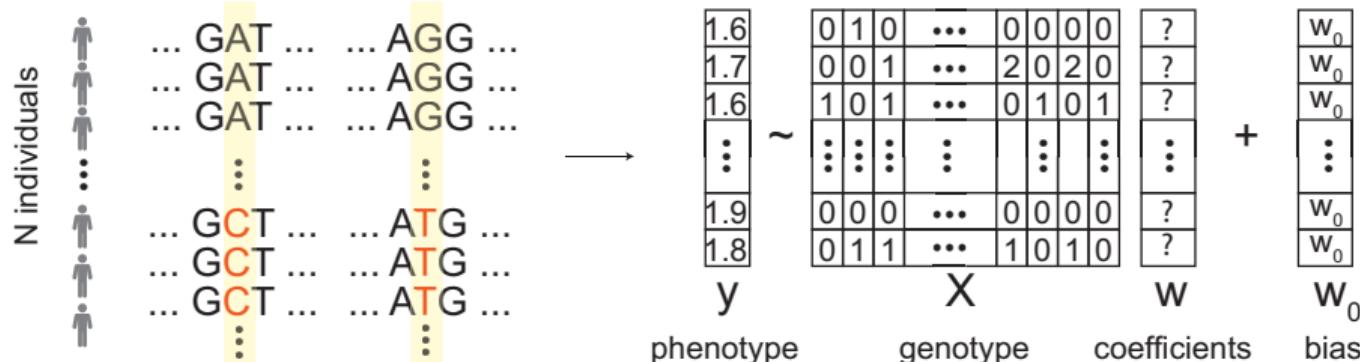
Non-linear basis function

Summary

Multiple linear regression

- Goal: learn how to fit a *multiple regression* to predict the outcome variable y using *multiple* input features
- We can write the response as a weighted linear sum of the input features:

$$f(\mathbf{x}^{(n)}; \mathbf{w}) = w_0 + w_1 x_1^{(n)} + w_2 x_2^{(n)} + \dots + w_D x_D^{(n)} = w_0 + \sum_{d=1}^D w_d x_d^{(n)}$$



Multiple regression in matrix form

Suppose N training examples and D features. The data matrices are:

- Response: $\mathbf{y} \in \mathbb{R}^{N \times 1}$
- Input feature matrix: $[\mathbf{1}, \mathbf{X}] \in \mathbb{R}^{N \times (D+1)}$
- Regression coefficients: $[w_0; \mathbf{w}] \in \mathbb{R}^{(D+1) \times 1}$

We can rewrite the multiple regression function as:

$$\begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(N)} \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \dots & x_D^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \dots & x_D^{(2)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(N)} & x_2^{(N)} & \dots & x_D^{(N)} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_D \end{bmatrix} + \begin{bmatrix} \epsilon^{(1)} \\ \epsilon^{(2)} \\ \vdots \\ \epsilon^{(N)} \end{bmatrix}$$

$$\mathbf{y} = \mathbf{1}w_0 + \mathbf{X}\mathbf{w} + \boldsymbol{\epsilon}$$

Our goal is to find the ordinary least square (OLS) solution for the coefficients \mathbf{w} :

$$w_0^*, \mathbf{w}^* \leftarrow \arg \min_{w_0, \mathbf{w}} \|\mathbf{y} - \mathbf{X}\mathbf{w} - w_0\|_2^2 = (\mathbf{y} - \mathbf{X}\mathbf{w} - w_0)^T(\mathbf{y} - \mathbf{X}\mathbf{w} - w_0)$$

Notation: $\|\mathbf{a}\|_2 = \sqrt{\sum_i a_i^2}$ is L_2 -norm and $\|\mathbf{a}\|_2^2 = \sum_i a_i^2$ is the squared of the L2-norm.

Multiple regression OLS derivation for the bias term w_0

Let the loss function be $J(\mathbf{w}, w_0) = (\mathbf{y} - \mathbf{X}\mathbf{w} - w_0)^T(\mathbf{y} - \mathbf{X}\mathbf{w} - w_0)$.

$$\begin{aligned}\frac{\partial J(\mathbf{w})}{\partial w_0} &= \frac{\partial}{\partial w_0} ((\mathbf{y} - \mathbf{X}\mathbf{w}) - \mathbf{1}w_0)^T((\mathbf{y} - \mathbf{X}\mathbf{w}) - \mathbf{1}w_0) \\ &= \frac{\partial}{\partial w_0} \{(\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w}) - (\mathbf{y} - \mathbf{X}\mathbf{w})^T\mathbf{1}w_0 - w_0\mathbf{1}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) + \mathbf{1}^T\mathbf{1}w_0^2\} \\ &= \frac{\partial}{\partial w_0}(\mathbf{y} - \mathbf{X}\mathbf{w})^T(\mathbf{y} - \mathbf{X}\mathbf{w}) - \frac{\partial}{\partial w_0}2w_0\mathbf{1}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) + \frac{\partial}{\partial w_0}Nw_0^2 \\ &= 0 - 2\mathbf{1}^T(\mathbf{y} - \mathbf{X}\mathbf{w}) + 2Nw_0 \\ &= -2\mathbf{1}^T\mathbf{y} + 2\mathbf{1}^T\mathbf{X}\mathbf{w} + 2Nw_0 \\ &= -2 \sum_n y^{(n)} + 2 \left(\sum_n \mathbf{x}^{(n)} \right) \mathbf{w} + 2Nw_0 \stackrel{\text{set}}{=} 0\end{aligned}$$

Solving $\frac{\partial J(\mathbf{w})}{\partial w_0} = 0$ for w_0 :

$$\hat{w}_0 = \frac{1}{N} \sum_n y^{(n)} - \frac{1}{N} \left(\sum_n \mathbf{x}^{(n)} \right) \mathbf{w} \equiv \bar{y} - \bar{\mathbf{x}}\mathbf{w}$$

Plugging in the OLS estimate for the bias term w_0 into the loss

$$\begin{aligned} J(\mathbf{w}, \hat{w}_0) &= \|\mathbf{y} - \mathbf{X}\mathbf{w} - \mathbf{1}\hat{w}_0\|_2^2 \\ &= \|\mathbf{y} - \mathbf{X}\mathbf{w} - \mathbf{1}(\bar{y} - \bar{\mathbf{x}}\mathbf{w})\|_2^2 \\ &= \|(\mathbf{y} - \mathbf{1}\bar{y}) - (\mathbf{X}\mathbf{w} - \mathbf{1}\bar{\mathbf{x}}\mathbf{w})\|_2^2 \\ &= \|(\mathbf{y} - \mathbf{1}\bar{y}) - (\mathbf{X} - \mathbf{1}\bar{\mathbf{x}})\mathbf{w}\|_2^2 \\ &\equiv \|\dot{\mathbf{y}} - \dot{\mathbf{X}}\mathbf{w}\|_2^2 \end{aligned}$$

where both the response and input features are mean-centered such that for each example n :

$$\begin{aligned} \dot{y}^{(n)} &= y^{(n)} - \bar{y} \\ \dot{\mathbf{x}}^{(n)} &= \mathbf{x}^{(n)} - \bar{\mathbf{x}} \end{aligned}$$

Note: the second equation is element-wise subtraction: for each feature d ,

$$\dot{x}_d^{(n)} = x_d^{(n)} - \bar{x}_d$$

Solving all multiple regression coefficients simultaneously

Here we can assume that either the data are centered (i.e., $\mathbf{y} = \hat{\mathbf{y}}$, $\mathbf{X} = \hat{\mathbf{X}}$) or we have $\mathbf{x}_0 = \mathbf{1}$. In the latter case, the estimate for w_0 is the first element in $\hat{\mathbf{w}}$.

$$\begin{aligned}\frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w}) = \frac{\partial}{\partial \mathbf{w}} (\mathbf{y}^\top \mathbf{y} - \underbrace{\mathbf{y}^\top \mathbf{X}\mathbf{w} - \mathbf{w}^\top \mathbf{X}^\top \mathbf{y}}_{2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y}} + \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w}) \\ &= \frac{\partial}{\partial \mathbf{w}} (\mathbf{y}^\top \mathbf{y} - 2\mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w}) = \underbrace{\frac{\partial}{\partial \mathbf{w}} \mathbf{y}^\top \mathbf{y}}_0 - 2 \frac{\partial}{\partial \mathbf{w}} \mathbf{w}^\top \mathbf{X}^\top \mathbf{y} + \frac{\partial}{\partial \mathbf{w}} \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w} \\ &\stackrel{\dagger}{=} -2\mathbf{X}^\top \mathbf{y} + 2\mathbf{X}^\top \mathbf{X}\mathbf{w} = 2\mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = 2\mathbf{X}^\top (\hat{\mathbf{y}} - \mathbf{y})\end{aligned}$$

[†]To get this equality, we make use of two general properties in matrix differentiation¹:

$$\frac{\partial \mathbf{b}^\top \mathbf{a}}{\partial \mathbf{b}} = \mathbf{a}, \quad \frac{\partial \mathbf{b}^\top \mathbf{A}\mathbf{b}}{\partial \mathbf{b}} = 2\mathbf{A}\mathbf{b}$$

Setting the derivative to zero and solve for \mathbf{w} gives the closed-form solution:

$$0 = -2\mathbf{X}^\top \mathbf{y} + 2\mathbf{X}^\top \mathbf{X}\mathbf{w} \rightarrow \mathbf{X}^\top \mathbf{X}\mathbf{w} = \mathbf{X}^\top \mathbf{y} \rightarrow \mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

¹Matrix cookbook Section 2.4 Eq 69 & Eq 85

Uniqueness of the OLS solution

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} \quad (9)$$

- The OLS solution is available when the $D \times D$ square matrix $\mathbf{A} = \mathbf{X}^\top \mathbf{X}$ is invertible.
- Matrix inverse \mathbf{A}^{-1} can be computed by eigendecomposition:

$$\mathbf{A} = \mathbf{Q} \Lambda \mathbf{Q}^{-1} \implies \mathbf{A}^{-1} = \mathbf{Q}^{-1} \Lambda^{-1} \mathbf{Q}$$

where

- \mathbf{Q} is the square $D \times D$ matrix, whose i^{th} column is the i^{th} eigenvector
- \mathbf{Q} is also an **orthonormal** matrix, which means $\mathbf{Q}^\top \mathbf{Q} = \mathbf{I} \implies \mathbf{Q}^\top = \mathbf{Q}^{-1}$
- Each eigenvector \mathbf{u} in \mathbf{Q} is orthonormal to itself: $\mathbf{u}^\top \mathbf{u} = 1$ and two different eigenvectors are orthogonal to each other: $\mathbf{u}^\top \mathbf{v} = 0$
- Λ is the diagonal matrix whose diagonal elements are the corresponding eigenvalues, i.e., $\Lambda_{i,i} = \lambda_i$
- Therefore, \mathbf{A} is not invertible if some of its eigenvalues are zeros, which can happen when two features are perfectly correlated, e.g., $\mathbf{x}_2 = 1 - \mathbf{x}_1$, meaning that the number of linearly independent columns (i.e., rank) is smaller than D .

Time complexity

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

- The inner product $\mathbf{A} = \mathbf{X}^\top \mathbf{X}$ takes $O(ND^2)$
- The inversion of the $D \times D$ matrix \mathbf{A}^{-1} takes $O(D^3)$
- Computing $\mathbf{X}^\top \mathbf{y}$ takes $O(ND)$

Therefore, the total time complexity is $O(ND^2 + D^3)$.

This can be very expensive or infeasible for large D (e.g., 1 million SNPs or 20,000 genes) and/or large N (e.g., half million individuals). For this reason, although we can derive closed-form solution, *Stochastic Gradient Descent* (SGD) is used for large dataset (Module 4.4).

Multivariate regression

We can also adapt the equation for multiple response variables. Instead of a response vector $\mathbf{y} \in \mathbb{R}^N$, we have a response matrix $\mathbf{Y} \in \mathbb{R}^{N \times C}$ for C response variables. The multivariate regression function is:

$$\mathbf{Y} = \mathbf{X}\mathbf{W} \tag{10}$$

where $\mathbf{W} \in \mathbb{R}^{D \times C}$.

The OLS solution for \mathbf{W} is then:

$$\mathbf{W} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y} \tag{11}$$

Note here the OLS coefficient \mathbf{w}_k for each response variable k is computed *independently* in the above solution. Therefore, the resulting OLS \mathbf{W} is identical to fitting each $D \times 1$ regression coefficient \mathbf{w}_k separately and then concatenate the C vectors together to form the matrix \mathbf{W} .

Outline

Objectives

Simple linear regression

Multiple linear regression

Probabilistic interpretation

Non-linear basis function

Summary

Probabilistic interpretation: Gaussian response variable

The general multivariate normal (MVN) distribution is:

$$p(\mathbf{y}|\mathbf{X}) = \mathcal{N}(\mathbf{y}|\mathbf{X}\mathbf{w}, \Sigma) = \det(2\pi\Sigma)^{-1/2} \exp\left(-\frac{1}{2}(\mathbf{y} - \mathbf{X}\mathbf{w})^\top \Sigma^{-1} (\mathbf{y} - \mathbf{X}\mathbf{w})\right) \quad (12)$$

where Σ is a $N \times N$ covariance matrix between the N samples. If we assume the samples are *i.i.d.*, then Σ is a diagonal matrix $\sigma^2 \mathbf{I}$, where \mathbf{I} is an identity matrix:

$$\Sigma \stackrel{i.i.d.}{=} \begin{bmatrix} \sigma^2 & 0 & 0 & \dots & 0 \\ 0 & \sigma^2 & 0 & \dots & 0 \\ 0 & 0 & \sigma^2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & \sigma^2 \end{bmatrix} = \sigma^2 \mathbf{I}, \quad \Sigma^{-1} = \sigma^{-2} \mathbf{I}, \quad \det(2\pi\Sigma)^{-1/2} = (2\pi\sigma^2)^{-N/2}$$

where $\det(\mathbf{A})$ is the determinant of a square matrix.

The MVN can then be simplified as the product of individual Gaussians:

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}) &= (2\pi\sigma^2)^{-N/2} \exp\left(-\frac{1}{2\sigma^2}(\mathbf{y} - \mathbf{X}\mathbf{w})^\top(\mathbf{y} - \mathbf{X}\mathbf{w})\right) \\ &= (2\pi\sigma^2)^{-N/2} \exp\left(-\frac{1}{2\sigma^2} \sum_n (y^{(n)} - \mathbf{x}^{(n)}\mathbf{w})^2\right) \end{aligned}$$

Taking the logarithm of the likelihood, we have

$$\begin{aligned} \ln p(\mathbf{y}|\mathbf{X}) &= \underbrace{-\frac{N}{2} \ln(2\pi\sigma^2)}_{\text{constant w.r.t. } \mathbf{w}} - \sum_n \left(\frac{1}{2\sigma^2} (y^{(n)} - \mathbf{x}^{(n)}\mathbf{w})^2 \right) \\ &\propto -\frac{1}{2\sigma^2} \sum_n (y^{(n)} - \mathbf{x}^{(n)}\mathbf{w})^2 \\ &= -\frac{1}{2\sigma^2} \|\mathbf{y} - \mathbf{X}\mathbf{w}\|_2^2 \end{aligned}$$

The last equation indicates that the log likelihood is proportional to the negative SSE.

Maximum likelihood estimation

Given that,

$$\ln p(\mathbf{y}|\mathbf{X}) \propto -\frac{1}{2\sigma^2} \sum_n (y^{(n)} - \mathbf{x}^{(n)} \mathbf{w})^2 = -\frac{1}{2\sigma^2} J(\mathbf{w})$$

Since σ^2 is a constant, minimizing SSE w.r.t. \mathbf{w} is equivalent to maximizing the Gaussian likelihood w.r.t. \mathbf{w} :

$$\arg \min_{\mathbf{w}} J(\mathbf{w}) = \arg \max_{\mathbf{w}} \ln p(\mathbf{y}|\mathbf{X})$$

The maximum likelihood estimator for \mathbf{w} is then identical to the OLS \mathbf{w} :

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

Outline

Objectives

Simple linear regression

Multiple linear regression

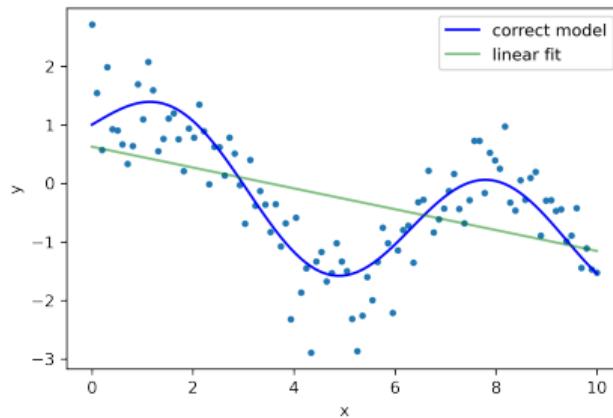
Probabilistic interpretation

Non-linear basis function

Summary

Fitting non-linear data by transforming the input features

Consider the toy dataset below. It is obvious that our attempt to model y as a linear function of $\hat{y} = w_0 + xw_1$ would produce a bad fit.



Idea: we can create more features using the given features. For example, we can use a M-degree polynomial function and treat each power degree as a standalone feature:

$$\hat{y} = x^0 w_0 + x^1 w_1 + x^2 w_2 + \dots + x^M w_M = \sum_{m=0}^M x^m w_m$$

Fitting non-linear data by transforming the input features

In general, we can transform input feature x with a (non-linear) **basis function** $\phi(x)$.
The multiple linear regression operates on the basis-transformed features:

$$\hat{y} = \sum_{m=0}^M \phi_m(x) w_m = \Phi(x)\mathbf{w} \quad (13)$$

We then simply replace all of the occurrences of x with $\Phi(x)$ in the OLS solution:

$$\hat{\mathbf{w}} = (\Phi(x)^T \Phi(x))^{-1} \Phi(x)^T \mathbf{y}, \quad \text{where}$$

$$\Phi(\mathbf{x}) = \begin{bmatrix} \phi_0(x^{(1)}) & \phi_1(x^{(1)}) & \dots & \phi_M(x^{(1)}) \\ \phi_0(x^{(2)}) & \phi_1(x^{(2)}) & \dots & \phi_M(x^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(x^{(N)}) & \phi_1(x^{(N)}) & \dots & \phi_M(x^{(N)}) \end{bmatrix}$$

Transforming high-dimensional features

- This can also be done on high-dimensional input feature $\mathbf{x}^{(n)}$ by transforming each feature with $x_d^{(n)}$ with say M -degree polynomial:

$$\Phi(x_d^{(n)}) = [(x_d^{(n)})^1, \dots, (x_d^{(n)})^M]$$

- We can then concatenate all transformed $D \times M$ features

$$\mathbf{x}^{(n)} = [(x_1^{(n)})^1, \dots, (x_1^{(n)})^M, \dots, (x_D^{(n)})^1, \dots, (x_D^{(n)})^M] = [\Phi(x_1^{(n)}), \dots, \Phi(x_D^{(n)})]$$

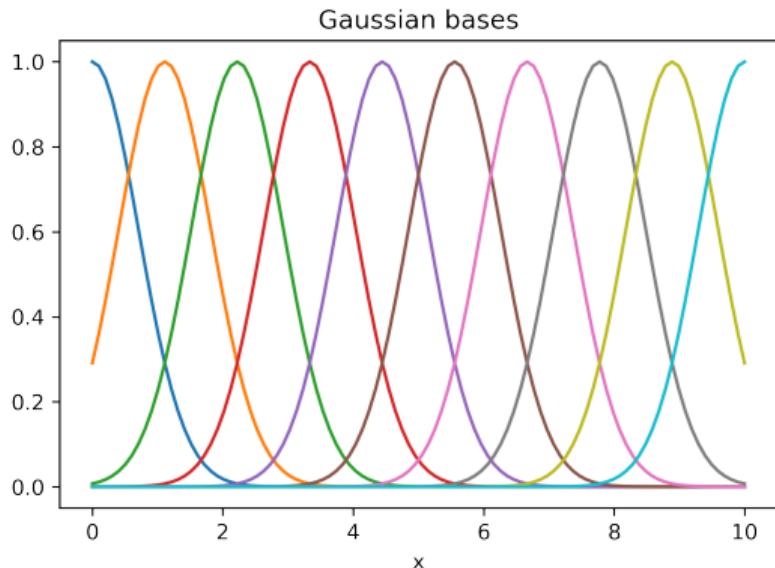
- This creates a input matrix of dimension $N \times (D \times M)$:

$$\Phi(\mathbf{X}) = \begin{bmatrix} \Phi(x_1^{(1)}) & \Phi(x_2^{(1)}) & \dots & \Phi(x_D^{(1)}) \\ \Phi(x_1^{(2)}) & \Phi(x_2^{(2)}) & \dots & \Phi(x_D^{(2)}) \\ \vdots & \vdots & \ddots & \vdots \\ \Phi(x_1^{(N)}) & \Phi(x_2^{(N)}) & \dots & \Phi(x_D^{(N)}) \end{bmatrix}$$

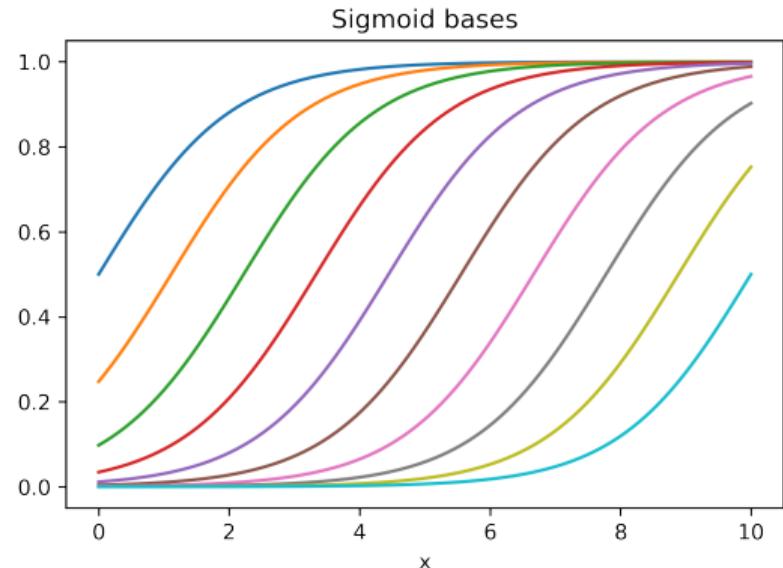
- Regression: $\hat{y} = w_0 + \Phi(\mathbf{X})\mathbf{w}$

Nonlinear basis functions

There are many nonlinear basis functions. Using scalar input $x \in \mathbb{R}$ as an example,



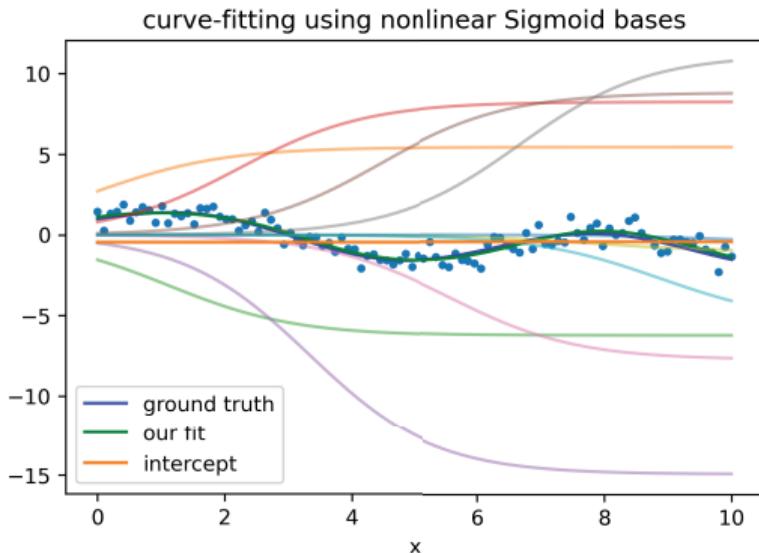
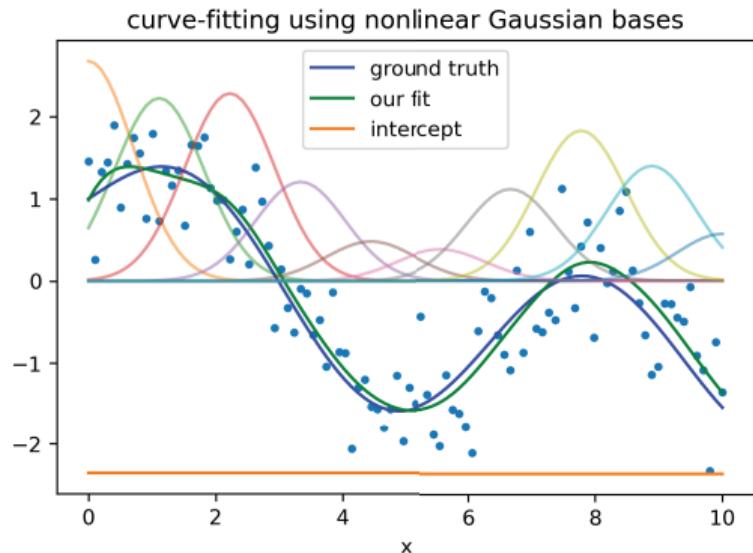
$$\phi_d(x) = \exp\left(-\frac{(x - \mu_d)^2}{2s^2}\right)$$



$$\phi_d(x) = \frac{1}{1 + \exp\left(-\frac{(x - \mu_d)^2}{s}\right)}$$

where each type of basis function has a different mean $\mu_d \in [0, 10]$ and $s = 1$.

Linear regression with nonlinear basis (See [Colab](#) for the implementations)



In both plots, the green curve (our fit) is the sum of weighted Gaussian bases (i.e., the colorful curves) plus the intercept:

$$\hat{y} = w_0 + \sum_d w_d \phi_d(\mathbf{x})$$

Summary

- Response variable y is modelled as a weighted linear sum of the features $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$
- We fit the model by minimizing the sum of squared errors (SSE)
$$\sum_n (y^{(n)} - \mathbf{X}^{(n)}\mathbf{w})^2$$
- OLS solution: $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ with $O(ND^2 + D^3)$ time complexity
- Minimizing SSE is equivalent to maximizing the log Gaussian likelihood given that the data points are *i.i.d.*
- Using non-linear basis functions to construct features can help model non-linear data. However, these non-linear basis functions are rigid and non-learnable.
- In Module 5.1, we will discuss neural networks which have *learnable* non-linear basis functions.